

Advanced probabilistic learning: Deep Belief Networks

Pietro Berkes, Brandeis University

Probabilistic learning

- an unsupervised model captures the distribution of the input data
- this distribution is usually best described in term of unobserved (hidden) causes



Generative models



objects




object parts




image features



pixels



Generative models and the brain




**our internal model:
a mirror version of the
real generative process?**

visual input
(light patterns)

external causes
(visual elements)

Unsupervised learning




**find hidden causes
that make input data
easy to describe**

visual input
(light patterns)

external causes
(visual elements)

Inference




**given an image, which
causes are present?**

visual input
(light patterns)


external causes
(visual elements)

Issues with hierarchical models

We would like to do this:



But we can only do this:



Issues with hierarchical models


- Many local minima: probability of observed data is very complicated function of parameters
- Some parameters are more sensitive than others
- In practice: global optimization of parameters is hell

Unrolling an RBM


Breakthrough in 2006: start by training a hierarchical network layer-by-layer

hidden
units, h


visible
units, v



Distribution over v can be found by Gibbs sampling:





Unrolling an RBM



Unrolling an RBM

The RBM was unrolled in a hierarchical generative model that defines the same distribution over the observed variables.

The activation functions and the weights must be the same everywhere!




Deep Belief Networks (2006)

- Hierarchical generative model
- Activation functions as in RBM:


$$a_i^{(l)} = b_i^{(l)} + \sum_j w_{ij}^{(l)} x_j^{(l+1)}$$

$$P(x_i^{(l)} = 1 \mid x^{(l+1)}) = \frac{1}{1 + e^{-a_i^{(l)}}}$$




Greedy learning

- Start assuming the weights are tied at all levels => the DBN is equivalent to an RBM
- We can train the DBN easily using the RBM training rule




Greedy learning

- Start assuming the weights are tied at all levels => the DBN is equivalent to an RBM
- We can train the DBN easily using the RBM training rule
- Fix the weights in the first layer, use the output of the first layer as the input for the next, and use the same trick




Greedy learning

- Start assuming the weights are tied at all levels => the DBN is equivalent to an RBM
- We can train the DBN easily using the RBM training rule
- Fix the weights in the first layer, use the output of the first layer as the input for the next, and use the same trick




Greedy learning

- Start assuming the weights are tied at all levels => the DBN is equivalent to an RBM
- We can train the DBN easily using the RBM training rule
- Fix the weights in the first layer, use the output of the first layer as the input for the next, and use the same trick
- And so on ...




Greedy learning

- Start assuming the weights are tied at all levels => the DBN is equivalent to an RBM
- We can train the DBN easily using the RBM training rule
- Fix the weights in the first layer, use the output of the first layer as the input for the next, and use the same trick
- And so on ...



Greedy learning

- Greedy learning or “learning by re-representing”
- Guaranteed to improve the model only until the second layer
- Works fairly well in practice
- After this greedy phase, one should proceed to a global fine-tuning of the weights



Demos

- Generating walks
- Generating digits

Hands-on

- I'll show how to write a DBN model for hand-written digits
- Greedy phase only
- Ingredients: RBM, RBM with labels

