

Exercises L3: Restricted Boltzmann Machine


Exercise 1

- Write a function `sample_A_or_B(prob)` that returns 'A' with probability `prob`, and 'B' with probability `(1.0 - prob)`.
(Use `numpy.random.rand` to generate random numbers.)
- Draw 1000 samples with `prob=0.3` and verify that you get approximately 300 A's and 700 B's

Exercise 2

Implement and test a function that samples from the hidden units of a RBM. Have a look at the files `sampling.py` and `test_sampling.py`.

- Consider a RBM with 2 visible units and 4 hidden units. The bias terms are set to zero, and the weights are as in the figure below.



Compute the probability of each hidden unit being 1, $P(h_j = 1 | v)$, for the following inputs:

$v = [v_0, v_1]$	[0, 0]	[1, 0]	[0, 1]	[1, 1]
h_0				
h_1				
h_2				
h_3				

- Complete the test in `test_sampling.py` with the values computed in a)
- Write the body of the function `sample_h` in `rbm.py` until the test passes

Exercise 3

The file `letter_S.npy` contains 39 handwritten examples of the letter “S”. We want to use a RBM to learn a model of this letter. We will use the RBM implementation in the machine learning library MDP.



The module `letter.py` defines two functions:

- `load_data`: load the data from the file `letter_S.npy` .
Each letter is a vector of length 320 (20x16 pixels)
- `show_letter`: display an activity pattern in a `matplotlib` figure

- a) Load the data and use the function `show_letter` to display the handwritten letters
- b) The object `mdp.nodes.RBMNode` is an implementation of a Restricted Boltzmann Machine algorithm. Read the documentation for the object: from the ipython prompt type

```
import mdp
mdp.nodes.RBMNode ?
mdp.nodes.RBMNode.train ?
mdp.nodes.RBMNode.sample_h ?
mdp.nodes.RBMNode.sample_v ?
```

- c) Create a RBM with 100 hidden units:
- ```
rbm = mdp.nodes.RBMNode(100)
```
- d) Train the network by presenting to it the data 100 times. Use `help(rbm.train)` to understand how to do that. You can keep the optional arguments at their default value or play around with them if you wish.
- e) During training, the network has learned a model of the data. We can use it to generate new examples of the letter “S”:

- a. start from a random input pattern:

```
v = np.random.randint(0, 2, size=(1, 320))
this is the probability of v=1
pv = v
```

- b. sample from the hidden layer given the input pattern, and again the visible layer given the new hidden pattern:

```
ph, h = rbm.sample_h(pv)
pv, v = rbm.sample_v(ph)
show(v)
```

Do it a few times until the RBM begins sampling from the equilibrium distribution. Each pattern is an exemplar of what the model has learned.

- c. plot the probability of the visible units being active instead:  
`show(pv)`